



Introduction to Distributed File Systems and HDFS



A Few Quotes to Start Us Off

Data is a precious thing and will last longer than the systems themselves.

(Tim Berners-Lee, 2008)

Data is not information, information is not knowledge, knowledge is not understanding, understanding is not wisdom.

(Sir Arthur Conan Doyle)



Learning Objectives

- **Describe basic file system concepts**
- **Discuss an overview of HDFS**
- **Use the HDFS CLI**
- **Use the HDFS API**
- **Exporting MapR-FS with NFS**




Describe Basic File System Concepts



Describe File System Concepts

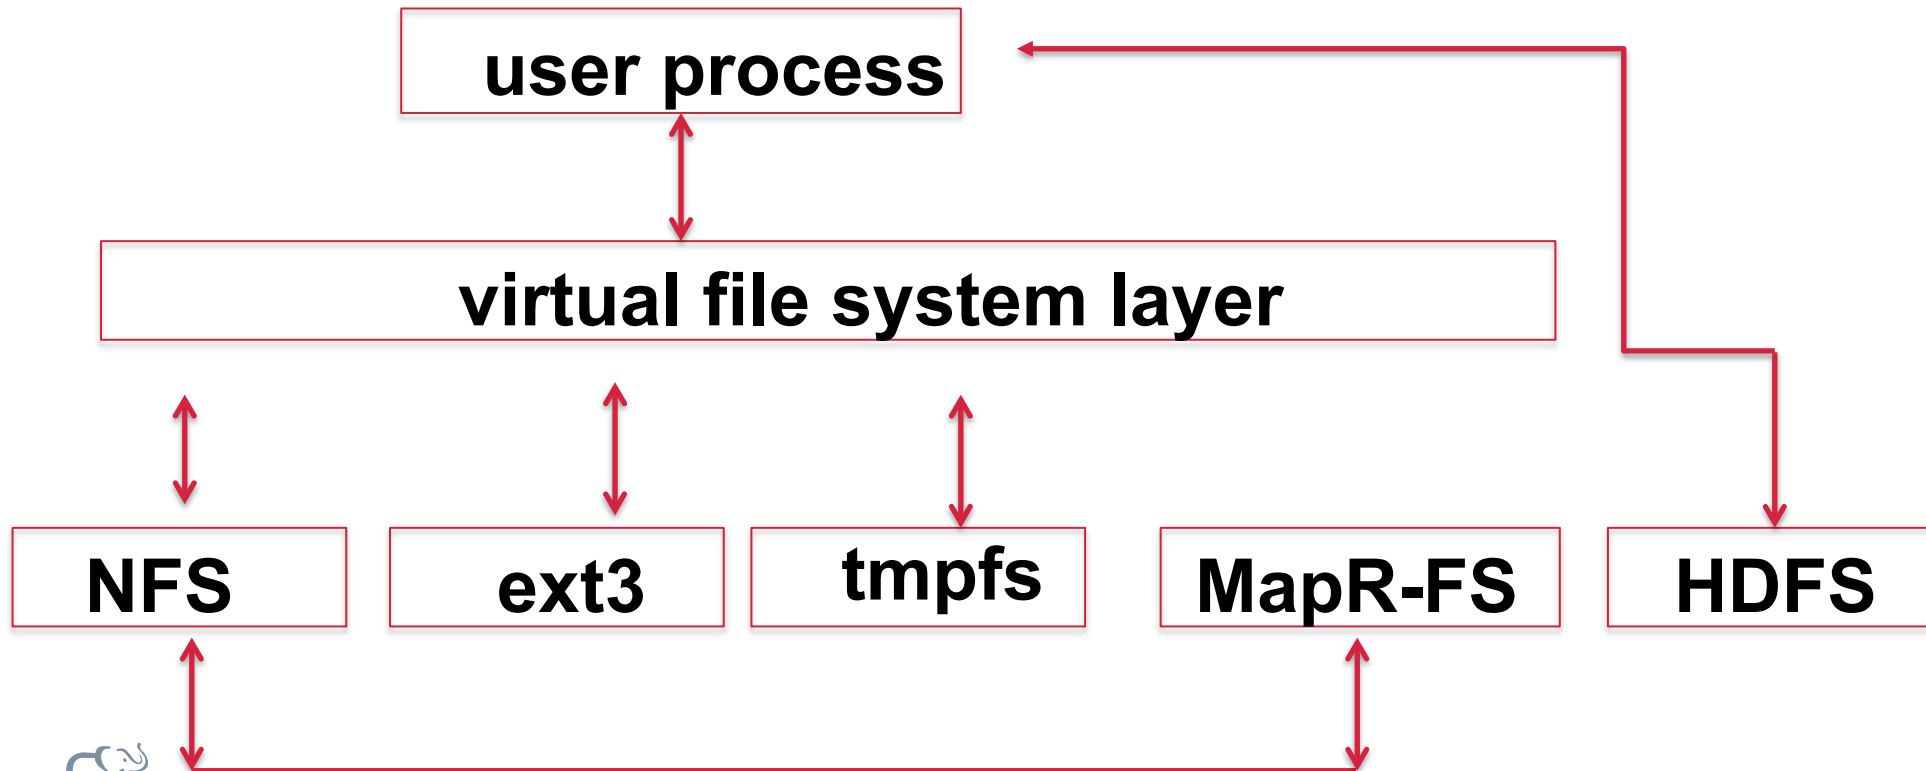
- Logical structure that organizes files on a storage medium
- Dictates how data is stored and retrieved
- Contains data and metadata



ext2 file system


Describe the Purpose of a Virtual File System (VFS)

- Translation layer from generic file system to real file system
- Enables standard POSIX file access



Describe Distributed File System Concepts


- Centrally stores metadata and distributes actual data
- Overcomes space, performance, and availability limitations of a single machine
- Abstracts data locality from client access



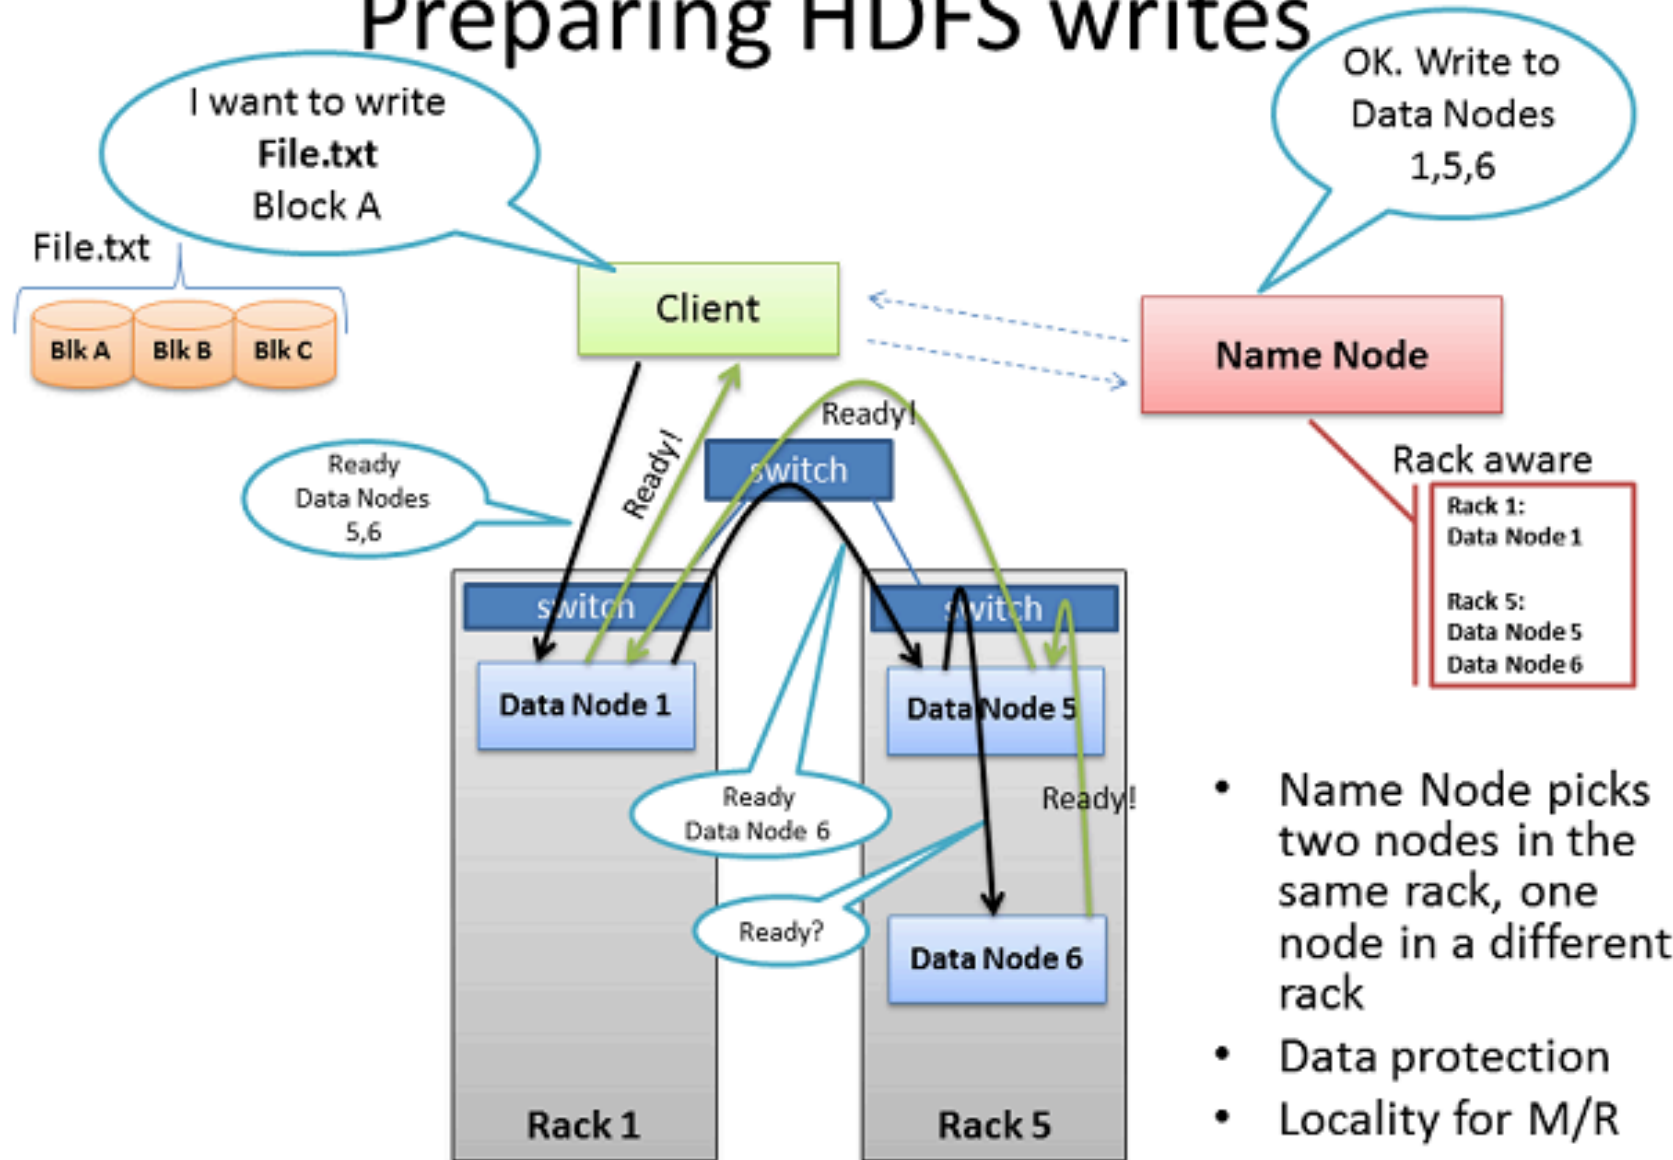
Discuss an Overview of HDFS



Describe the High-level HDFS Architecture



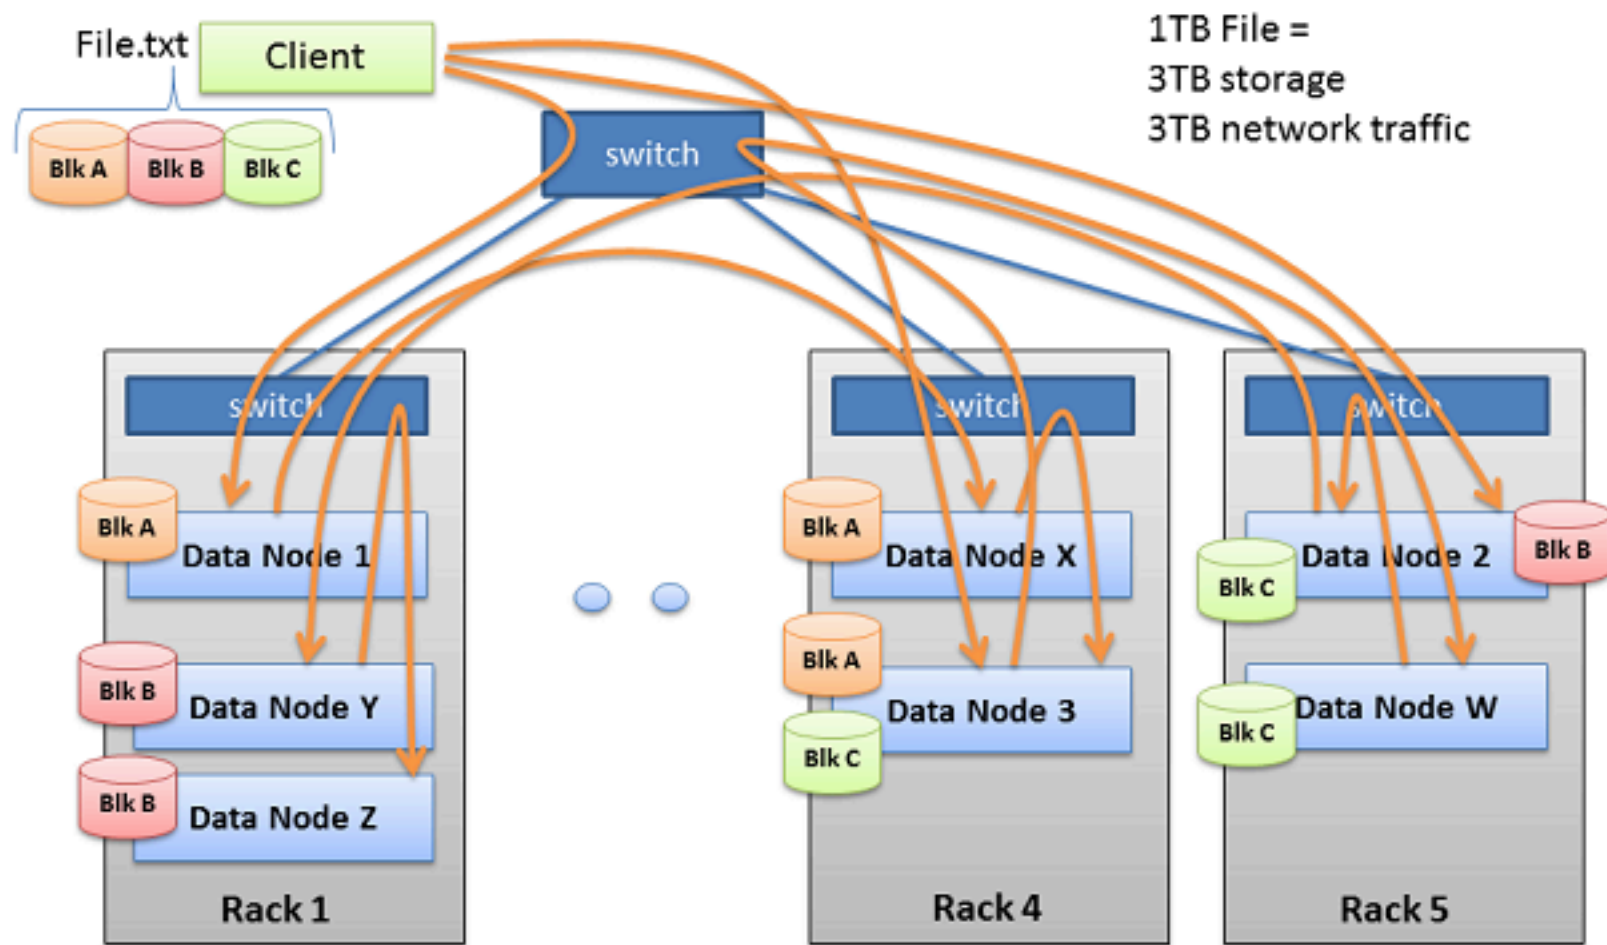
Preparing HDFS writes



- Name Node picks two nodes in the same rack, one node in a different rack
- Data protection
- Locality for M/R

BRAD HEDLUND .com

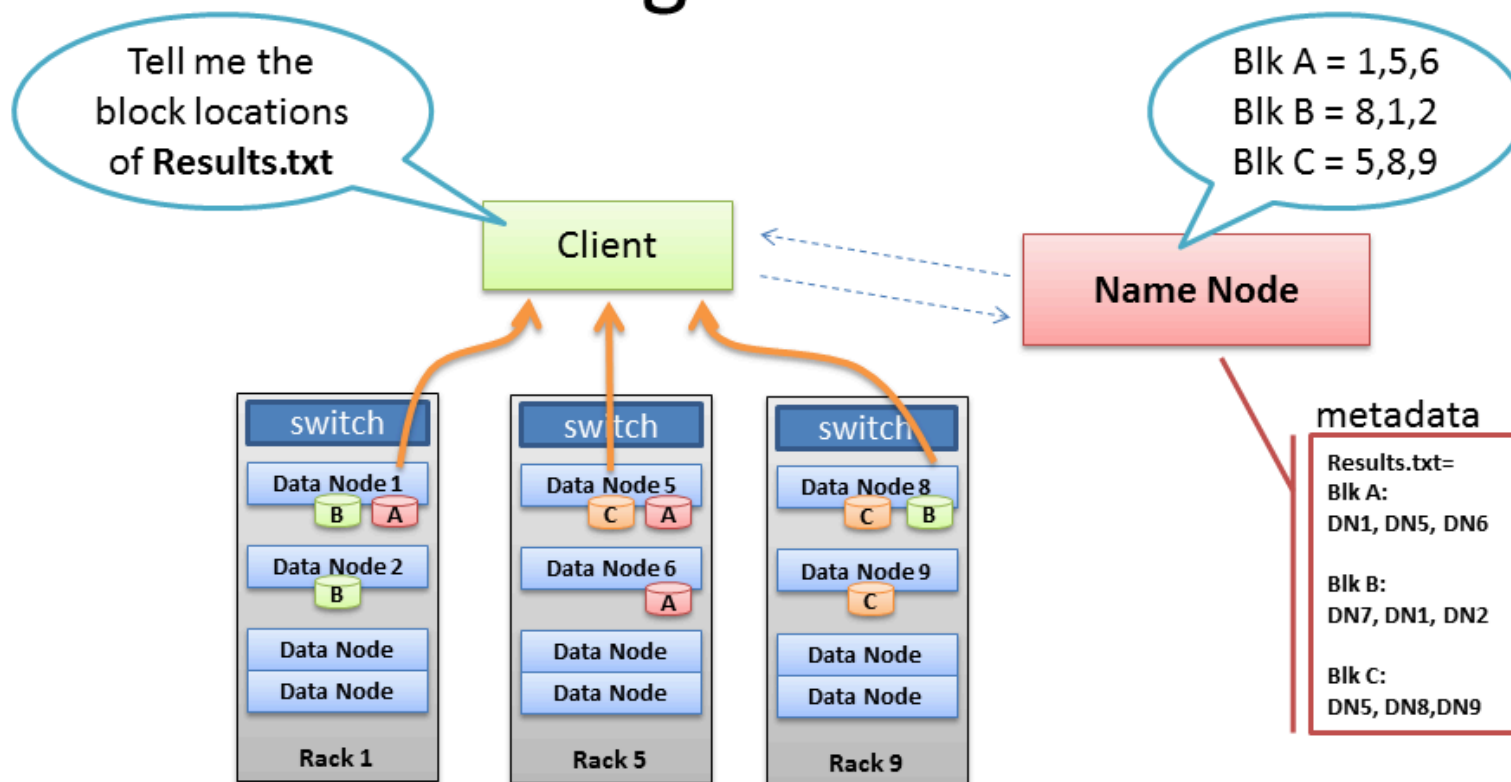
Multi-block Replication Pipeline



BRAD HEDLUND .com



Client reading files from HDFS



- Client receives Data Node list for each block
- Client picks first Data Node for each block
- Client reads blocks sequentially

BRAD HEDLUND .com



Identify Limitations of HDFS

aspect	limitation
Block size	Same size used for I/O, replication, and sharding
Mutability	Write-once, read-many (WORM)
POSIX semantics	Must use HDFS API/CLI to access data
Availability	No snapshot or built-in mirroring capability
Scalability	Namenode only scales to 100M files
Performance	Written in Java and runs on block device



Cite Differences b/w MapR-FS and HDFS

aspect	feature
Block size	Different sizes used for sharding, replicating, and performing I/O
Mutability	Full read-write capability
Access	Can NFS-mount MapR-FS volumes
POSIX semantics	Can use native OS to access data
Availability	Snapshots and local/remote mirroring support
Scalability	No limit to the number of files
Performance	Written in C and runs on raw device



Using the HDFS CLI



Use the `hadoop fs` CLI

Usage: `hadoop fs [command] [args]`

```
hadoop fs -mkdir mydir
```

```
hadoop fs -copyFromLocal /etc/hosts mydir
```

```
hadoop fs -lsr mydir
```

```
hadoop fs -cat mydir/hosts
```

```
hadoop fs -rm mydir/hosts
```



Differentiate Absolute and Relative Paths

```
$ hadoop fs -ls /  
data1 data2 tmp user var
```

```
$ hadoop conf -dump | grep fs.default.name  
fs.default.name=maprfs:///
```

```
$ hadoop fs -ls  
/user/jcasaletto/IN /user/jcasaletto/OUT
```

```
$ hadoop conf -dump | grep fs.mapr.working.dir  
fs.mapr.working.dir=/user/$USERNAME/
```



Use the `hadoop mfs` CLI

Usage: `hadoop mfs [command] [args]`

```
hadoop mfs -ln mydir yourdir
```

```
hadoop mfs -setcompression off mydir
```

```
hadoop mfs -setchunksize 65536 mydir
```

```
hadoop mfs -lnh original-file new-file
```



Use the Operating System CLI

```
mkdir /user/james/mydir
```

```
cp /etc/hosts /user/james/mydir
```

```
ls -R /user/james/mydir
```

```
ln -s /user/james/mydir /user/james/yourdir
```

```
rm /user/james/mydir/hosts
```

```
tail / grep / awk / sed
```



Using the HDFS Java API



Sample code

```
/* Copyright (c) 2009 & onwards. MapR Tech, Inc., All rights reserved */

//package com.mapr.fs;

import java.net.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.conf.*;

/**
 * Assumes mapr installed in /opt/mapr
 *
 * Compilation:
 * javac -cp $(hadoop classpath) MapRTest.java
 *
 * Run:
 * java -cp .:$(hadoop classpath) MapRTest /test
 */
```



Sample code (2)

```
public class MapRTest
{
    public static void main(String args[]) throws Exception {
        byte buf[] = new byte[ 65*1024];
        int ac = 0;
        if (args.length != 1) {
            System.out.println("usage: MapRTest pathname");
            return;
        }

        // maprfs:/// -> uses the first entry in /opt/mapr/conf/mapr-clusters.conf
        // maprfs:///mapr/my.cluster.com/
        // /mapr/my.cluster.com/

        // String uri = "maprfs:///";
        String dirname = args[ac++];

        Configuration conf = new Configuration();

        //FileSystem fs = FileSystem.get(URI.create(uri), conf); // if wanting to use a different cluster
        FileSystem fs = FileSystem.get(conf);
    }
}
```



Sample Code (3)

```
Path dirpath = new Path( dirname + "/dir" );
    Path wfilepath = new Path( dirname + "/file.w" );
    //Path rfilepath = new Path( dirname + "/file.r" );
    Path rfilepath = wfilepath;

    // try mkdir
    boolean res = fs.mkdirs( dirpath );
    if (!res) {
        System.out.println("mkdir failed, path: " +
dirpath);
    }
    return;
```



Sample code (4)

```
System.out.println( "mkdir( " + dirpath + ") went ok, now writing file");

// create wfile
FSDataOutputStream ostr = fs.create( wfilepath,
    true, // overwrite
    512, // buffersize
    (short) 1, // replication
    (long) (64*1024*1024) // chunksize
);
ostr.write(buf);
ostr.close();

System.out.println( "write( " + wfilepath + ") went ok");

// read rfile
System.out.println( "reading file: " + rfilepath);
FSDataInputStream istr = fs.open( rfilepath);
int bb = istr.readInt();
istr.close();
System.out.println( "Read ok");
}
}
```



Exporting MapR-FS with NFS

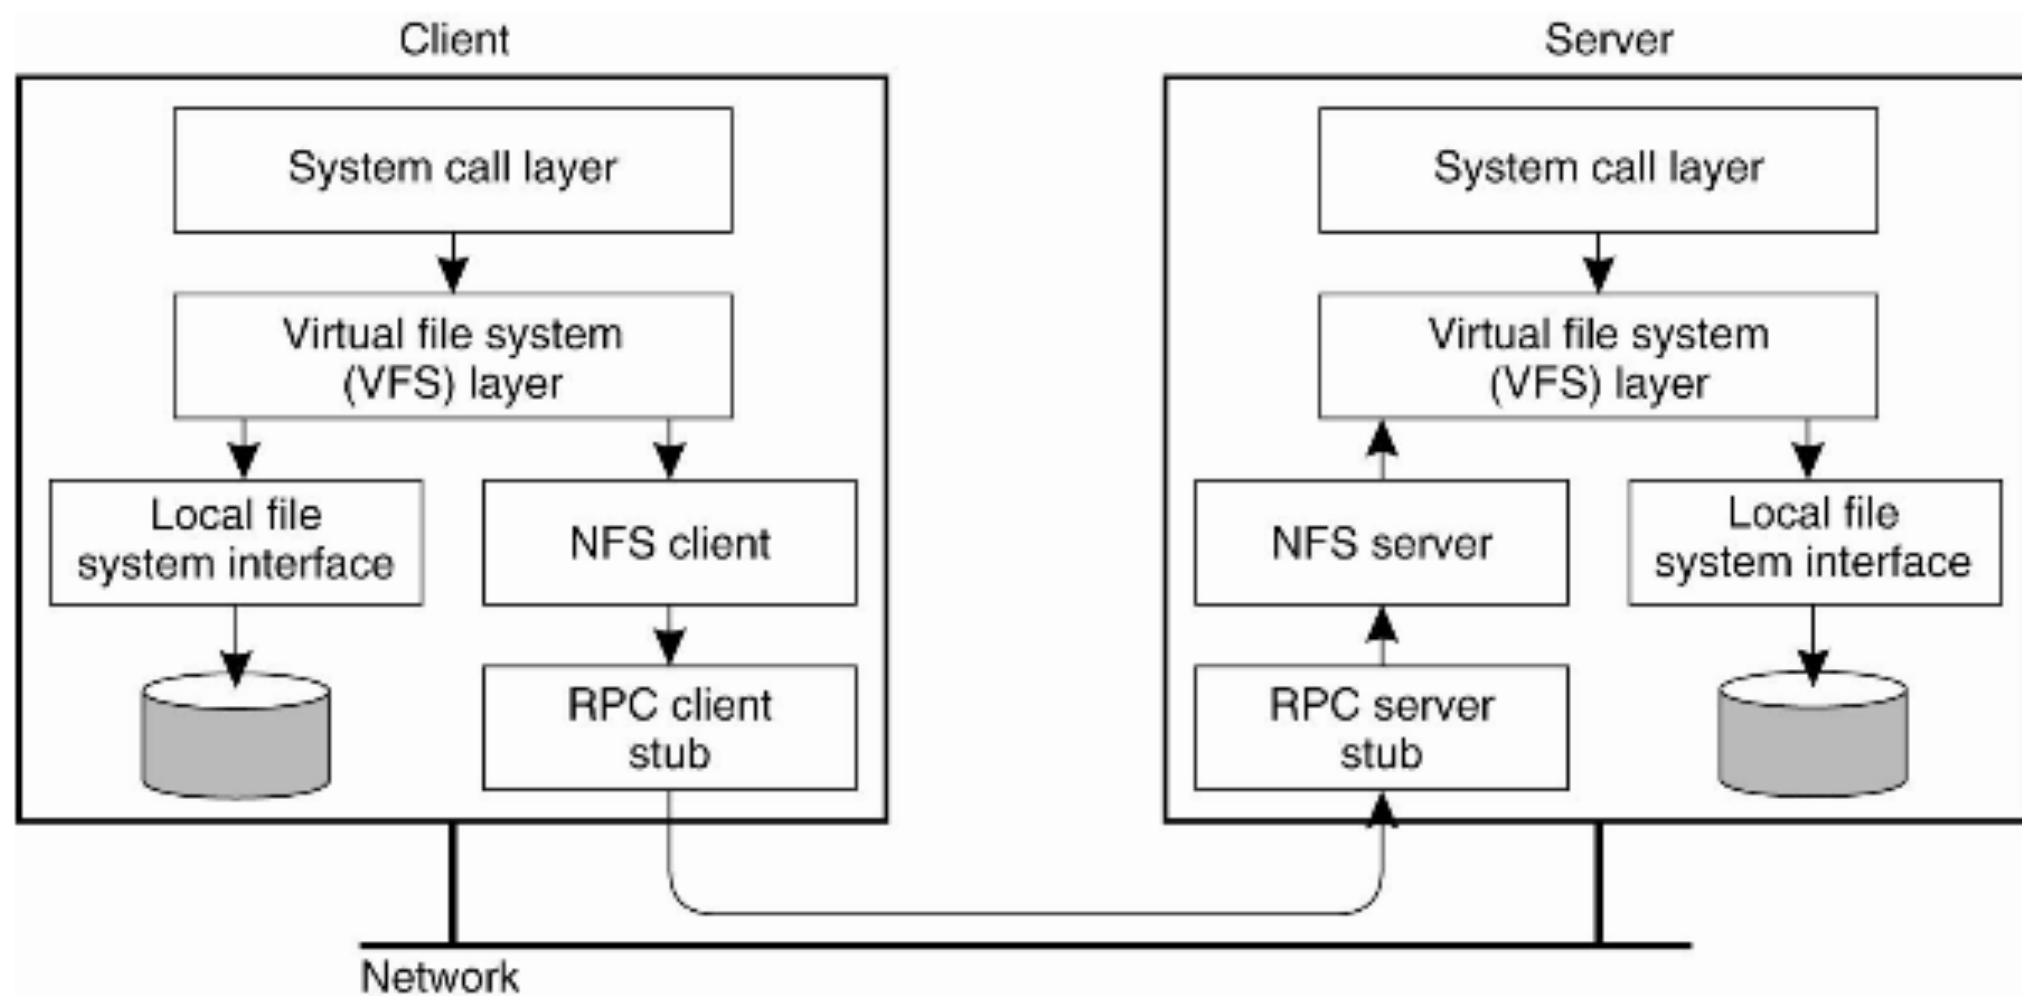


What is NFS?

- NFS = network file system
- Technology developed by Sun Microsystems in 1984
- Leverages VFS to present remote file system as if local
- NFS v1, v2, v3, v4 (MapR-FS supports v3)
- Runs as RPC service (MapR uses custom version of RPC)



NFS architecture



How to configure an NFS service

- **In Linux:**

- edit /etc/exports:

```
dir1 host1(option1, option2, ...) host2(option1, option2, ...)
dir2 host1(option1, option2, ...) host2(option1, option2, ...)
```

example:

```
/usr/local 192.168.0.1(ro) 192.168.0.2(ro)
/home      192.168.0.1(rw,nosuid) 192.168.0.2(rw,nosuid)
```

- **For MapR:**

- edit /opt/mapr/conf/exports



How to configure an NFS client

- **In standard Linux environment:**

- `edit /etc/fstab`

example:

```
venus.mapr.com:/user /mnt/user nfs nolock,actimeo=0,soft 0 0
```

- **In MapR environment:**

- `edit /opt/mapr/conf/mapr_fstab`

example:

```
localhost:/mapr/user /user nolock,soft,intr
```

